

Data Science Regression Project: Predicting Home Prices in Bangalore

Outline

1. Data Cleaning
2. Feature Engineering
3. Outlier Removal
4. Model Building

Dataset is downloaded from here: <https://www.kaggle.com/amitabhajoy/bengaluru-house-price-data> (<https://www.kaggle.com/amitabhajoy/bengaluru-house-price-data>)

```
In [2]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
import matplotlib
matplotlib.rcParams["figure.figsize"] = (20,10)
```

Data Load: Load banglore home prices into a dataframe

```
In [3]: df1 = pd.read_csv("bengaluru_house_prices.csv")
df1.head()
```

Out[3]:

| | area_type | availability | location | size | society | total_sqft | bath | balcony | price |
|---|---------------------|---------------|--------------------------|-----------|---------|------------|------|---------|--------|
| 0 | Super built-up Area | 19-Dec | Electronic City Phase II | 2 BHK | Coomee | 1056 | 2.0 | 1.0 | 39.07 |
| 1 | Plot Area | Ready To Move | Chikka Tirupathi | 4 Bedroom | Theanmp | 2600 | 5.0 | 3.0 | 120.00 |
| 2 | Built-up Area | Ready To Move | Uttarahalli | 3 BHK | NaN | 1440 | 2.0 | 3.0 | 62.00 |
| 3 | Super built-up Area | Ready To Move | Lingadheeranahalli | 3 BHK | Soiewre | 1521 | 3.0 | 1.0 | 95.00 |
| 4 | Super built-up Area | Ready To Move | Kothanur | 2 BHK | NaN | 1200 | 2.0 | 1.0 | 51.00 |

```
In [4]: df1.shape
```

```
Out[4]: (13320, 9)
```

```
In [5]: df1.columns
```

```
Out[5]: Index(['area_type', 'availability', 'location', 'size', 'society',  
              'total_sqft', 'bath', 'balcony', 'price'],  
              dtype='object')
```

```
In [6]: df1['area_type'].unique()
```

```
Out[6]: array(['Super built-up Area', 'Plot Area', 'Built-up Area',  
              'Carpet Area'], dtype=object)
```

```
In [7]: df1['area_type'].value_counts()
```

```
Out[7]: Super built-up Area    8790  
        Built-up Area         2418  
        Plot Area             2025  
        Carpet Area             87  
        Name: area_type, dtype: int64
```

Drop features that are not required to build our model

```
In [9]: df2 = df1.drop(['area_type', 'society', 'balcony', 'availability'], axis='columns')  
        df2.shape
```

```
Out[9]: (13320, 5)
```

1. Data Cleaning: Handle NA values

```
In [10]: df2.isnull().sum()
```

```
Out[10]: location      1  
        size          16  
        total_sqft     0  
        bath          73  
        price          0  
        dtype: int64
```

```
In [11]: df2.shape
```

```
Out[11]: (13320, 5)
```

```
In [12]: df3 = df2.dropna()
df3.isnull().sum()
```

```
Out[12]: location      0
size                0
total_sqft          0
bath                0
price              0
dtype: int64
```

```
In [13]: df3.shape
```

```
Out[13]: (13246, 5)
```

In size column we see 4 BHK and 4 bedrooms which are the same exist. By using unique method we can see all the unique values in the column.

```
In [15]: df3['size'].unique()
```

```
Out[15]: array(['2 BHK', '4 Bedroom', '3 BHK', '4 BHK', '6 Bedroom', '3 Bedroom',
                '1 BHK', '1 RK', '1 Bedroom', '8 Bedroom', '2 Bedroom',
                '7 Bedroom', '5 BHK', '7 BHK', '6 BHK', '5 Bedroom', '11 BHK',
                '9 BHK', '9 Bedroom', '27 BHK', '10 Bedroom', '11 Bedroom',
                '10 BHK', '19 BHK', '16 BHK', '43 Bedroom', '14 BHK', '8 BHK',
                '12 Bedroom', '13 BHK', '18 Bedroom'], dtype=object)
```

We want to create a new column and get the first element of size as an integer number with following procedure:

1. Apply a function on size column by using lambda
2. Split the first element on size column
3. Convert it into integer.

Add new feature(integer) for bhk (Bedrooms Hall Kitchen)

```
In [16]: df3['bhk'] = df3['size'].apply(lambda x: int(x.split(' ')[0]))
df3.bhk.unique()
```

```
<ipython-input-16-681cf3aca53d>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df3['bhk'] = df3['size'].apply(lambda x: int(x.split(' ')[0]))
```

```
Out[16]: array([ 2,  4,  3,  6,  1,  8,  7,  5, 11,  9, 27, 10, 19, 16, 43, 14, 12,
                13, 18], dtype=int64)
```

I want to check the bhk>20

```
In [17]: df3[df3.bhk>20]
```

Out[17]:

| | location | size | total_sqft | bath | price | bhk |
|------|---------------------------|------------|------------|------|-------|-----|
| 1718 | 2Electronic City Phase II | 27 BHK | 8000 | 27.0 | 230.0 | 27 |
| 4684 | Munnekollal | 43 Bedroom | 2400 | 40.0 | 660.0 | 43 |

43 Bedrooms?!?=> This looks like an error with 2400 total_sqft

I want to check the values on total_sqft

```
In [18]: df3['total_sqft'].unique()
```

```
Out[18]: array(['1056', '2600', '1440', ..., '1133 - 1384', '774', '4689'],  
              dtype=object)
```

The value **1133-1384** is not a single number it is a range & I want to convert it into single value. One way is to do that I can take an average of two numbers.

First we want to see what kind of variations we have. In order to do that we can detect the number in total_sqft is float or not.

We will define a function is called is_float. The way this function work I will try to convert each value into float and if it is not a valid value such as this range (**1133-1384**) it will come into this except block and it returns False.

Explore total_sqft feature

```
In [19]: def is_float(x):  
          try:  
              float(x)  
          except:  
              return False  
          return True
```

Now we can apply is_float function into df3['total_sqft']. In order to see the value that it is not a float number we can do ~ operation.

```
In [20]: df3[~df3['total_sqft'].apply(is_float)].head(10)
```

Out[20]:

| | location | size | total_sqft | bath | price | bhk |
|-----|--------------------|-----------|----------------|------|---------|-----|
| 30 | Yelahanka | 4 BHK | 2100 - 2850 | 4.0 | 186.000 | 4 |
| 122 | Hebbal | 4 BHK | 3067 - 8156 | 4.0 | 477.000 | 4 |
| 137 | 8th Phase JP Nagar | 2 BHK | 1042 - 1105 | 2.0 | 54.005 | 2 |
| 165 | Sarjapur | 2 BHK | 1145 - 1340 | 2.0 | 43.490 | 2 |
| 188 | KR Puram | 2 BHK | 1015 - 1540 | 2.0 | 56.800 | 2 |
| 410 | Kengeri | 1 BHK | 34.46Sq. Meter | 1.0 | 18.500 | 1 |
| 549 | Hennur Road | 2 BHK | 1195 - 1440 | 2.0 | 63.770 | 2 |
| 648 | Arekere | 9 Bedroom | 4125Perch | 9.0 | 265.000 | 9 |
| 661 | Yelahanka | 2 BHK | 1120 - 1145 | 2.0 | 48.130 | 2 |
| 672 | Bettahalsoor | 4 Bedroom | 3090 - 5002 | 4.0 | 445.000 | 4 |

Above shows that total_sqft can be a range (e.g. 2100-2850). For such case we can just take average of min and max value in the range. There are other cases such as 34.46Sq. Meter which one can convert to square ft using unit conversion. I am going to just drop such corner cases to keep things simple

In order to clean this data we want to take an average value of the ranges one. Also, about some data like 34.46Sq. Mete or 4125Perch you can change the unit from Perch to sqft but here for simplicity we just ignore that data.

```
In [21]: def convert_sqft_to_num(x):
          tokens = x.split('-')
          if len(tokens) == 2:
              return (float(tokens[0])+float(tokens[1]))/2
          try:
              return float(x)
          except:
              return None
```

```
In [22]: df4 = df3.copy()
df4.total_sqft = df4.total_sqft.apply(convert_sqft_to_num)
df4 = df4[df4.total_sqft.notnull()]
df4.head(5)
```

Out[22]:

| | location | size | total_sqft | bath | price | bhk |
|---|--------------------------|-----------|------------|------|--------|-----|
| 0 | Electronic City Phase II | 2 BHK | 1056.0 | 2.0 | 39.07 | 2 |
| 1 | Chikka Tirupathi | 4 Bedroom | 2600.0 | 5.0 | 120.00 | 4 |
| 2 | Uttarahalli | 3 BHK | 1440.0 | 2.0 | 62.00 | 3 |
| 3 | Lingadheeranahalli | 3 BHK | 1521.0 | 3.0 | 95.00 | 3 |
| 4 | Kothanur | 2 BHK | 1200.0 | 2.0 | 51.00 | 2 |

For below row, it shows total_sqft as 2475 which is an average of the range 2100-2850

```
In [23]: df4.loc[30]
```

```
Out[23]: location      Yelahanka
size              4 BHK
total_sqft        2475
bath              4
price             186
bhk              4
Name: 30, dtype: object
```

```
In [24]: (2100+2850)/2
```

Out[24]: 2475.0

2. Feature Engineering

```
In [25]: df5 = df4.copy()
df5.head()
```

Out[25]:

| | location | size | total_sqft | bath | price | bhk |
|---|--------------------------|-----------|------------|------|--------|-----|
| 0 | Electronic City Phase II | 2 BHK | 1056.0 | 2.0 | 39.07 | 2 |
| 1 | Chikka Tirupathi | 4 Bedroom | 2600.0 | 5.0 | 120.00 | 4 |
| 2 | Uttarahalli | 3 BHK | 1440.0 | 2.0 | 62.00 | 3 |
| 3 | Lingadheeranahalli | 3 BHK | 1521.0 | 3.0 | 95.00 | 3 |
| 4 | Kothanur | 2 BHK | 1200.0 | 2.0 | 51.00 | 2 |

we want to create a new column which is called price_per_sqft. This value is really important in price market. This feature will help us do some outlier cleaning in the later stage. So we are doing some feature engineering which can be helpful for outlier detection and removal.

Add new feature called price per square feet

```
In [26]: df5['price_per_sqft'] = df5['price']*100000/df5['total_sqft']
df5.head()
```

Out[26]:

| | location | size | total_sqft | bath | price | bhk | price_per_sqft |
|---|--------------------------|-----------|------------|------|--------|-----|----------------|
| 0 | Electronic City Phase II | 2 BHK | 1056.0 | 2.0 | 39.07 | 2 | 3699.810606 |
| 1 | Chikka Tirupathi | 4 Bedroom | 2600.0 | 5.0 | 120.00 | 4 | 4615.384615 |
| 2 | Uttarahalli | 3 BHK | 1440.0 | 2.0 | 62.00 | 3 | 4305.555556 |
| 3 | Lingadheeranahalli | 3 BHK | 1521.0 | 3.0 | 95.00 | 3 | 6245.890861 |
| 4 | Kothanur | 2 BHK | 1200.0 | 2.0 | 51.00 | 2 | 4250.000000 |

Because price is in lagropies and in order to convert it to dollar we have to multiply it by 100000.

```
In [29]: df5_stats = df5['price_per_sqft'].describe()
df5_stats
```

```
Out[29]: count    1.320000e+04
mean       7.920759e+03
std        1.067272e+05
min        2.678298e+02
25%        4.267701e+03
50%        5.438331e+03
75%        7.317073e+03
max        1.200000e+07
Name: price_per_sqft, dtype: float64
```

```
In [30]: df5.to_csv("bhp.csv",index=False)
```

We'll now explore location column. I want to check how many locations are there and how many rows are available in my dataset for location.

Location is a categorical feature and it is text data. If you have too many locations it creates problem.

```
In [27]: # First check how many locations
len(df5.location.unique())
```

Out[27]: 1298

The number 1304 is too big and you have too many features. There are techniques to reduce the dimensions. One of the effective techniques is to come up with **other category**. Other category means when you have 1304 locations you will find there are many locations which have only 1 or 2 datapoints.

```
In [28]: df5.location.unique()
```

```
Out[28]: array(['Electronic City Phase II', 'Chikka Tirupathi', 'Uttarahalli', ...,  
              '12th cross srinivas nagar banshankari 3rd stage',  
              'Havanur extension', 'Abshot Layout'], dtype=object)
```

Examine locations which is a categorical variable. We need to apply dimensionality reduction technique here to reduce number of locations

```
In [31]: # First: Remove space between or end of the characters by using strip  
df5.location = df5.location.apply(lambda x: x.strip())  
location_stats = df5['location'].value_counts(ascending=False)  
location_stats
```

```
Out[31]: Whitefield          533  
Sarjapur Road             392  
Electronic City           304  
Kanakpura Road            264  
Thanisandra               235  
...  
Ramamohanapuram           1  
Amco Colony               1  
Popular Colony            1  
mvj engineering college   1  
2Electronic City Phase II  1  
Name: location, Length: 1287, dtype: int64
```

```
In [32]: location_stats.values.sum()
```

```
Out[32]: 13200
```

```
In [33]: len(location_stats[location_stats>10])
```

```
Out[33]: 240
```

```
In [34]: len(location_stats)
```

```
Out[34]: 1287
```

```
In [35]: len(location_stats[location_stats<=10])
```

```
Out[35]: 1047
```

Dimensionality Reduction

Any location having less than 10 data points should be tagged as "other" location. This way number of categories can be reduced by huge amount. Later on when we do one hot encoding, it will help us with having fewer dummy columns


```
In [36]: location_stats_less_than_10 = location_stats[location_stats<=10]
location_stats_less_than_10
```

```
Out[36]: Basapura          10
Nagadevanahalli         10
Dairy Circle            10
Sadashiva Nagar         10
Thyagaraja Nagar        10
..
Ramamohanapuram         1
Amco Colony             1
Popular Colony          1
mvj engineering college 1
2Electronic City Phase II 1
Name: location, Length: 1047, dtype: int64
```

```
In [37]: len(df5.location.unique())
```

```
Out[37]: 1287
```

```
In [38]: df5.location = df5.location.apply(lambda x: 'other' if x in location_stats_less_t
len(df5.location.unique())
```

```
Out[38]: 241
```

By using this function any value <=10 is other and the others which are >10 keep as x. As you see the number is reduced from 1287 to 241. This is great because when I want to convert it to One Hot Encoding I will only have 283 DataFrames.

```
In [39]: df5.head(10)
```

```
Out[39]:
```

| | location | size | total_sqft | bath | price | bhk | price_per_sqft |
|---|--------------------------|-----------|------------|------|--------|-----|----------------|
| 0 | Electronic City Phase II | 2 BHK | 1056.0 | 2.0 | 39.07 | 2 | 3699.810606 |
| 1 | Chikka Tirupathi | 4 Bedroom | 2600.0 | 5.0 | 120.00 | 4 | 4615.384615 |
| 2 | Uttarahalli | 3 BHK | 1440.0 | 2.0 | 62.00 | 3 | 4305.555556 |
| 3 | Lingadheeranahalli | 3 BHK | 1521.0 | 3.0 | 95.00 | 3 | 6245.890861 |
| 4 | Kothanur | 2 BHK | 1200.0 | 2.0 | 51.00 | 2 | 4250.000000 |
| 5 | Whitefield | 2 BHK | 1170.0 | 2.0 | 38.00 | 2 | 3247.863248 |
| 6 | Old Airport Road | 4 BHK | 2732.0 | 4.0 | 204.00 | 4 | 7467.057101 |
| 7 | Rajaji Nagar | 4 BHK | 3300.0 | 4.0 | 600.00 | 4 | 18181.818182 |
| 8 | Marathahalli | 3 BHK | 1310.0 | 3.0 | 63.25 | 3 | 4828.244275 |
| 9 | other | 6 Bedroom | 1020.0 | 6.0 | 370.00 | 6 | 36274.509804 |

3. Outlier Removal Using Business Logic

Outliers are the data points which could be either data errors or represent the extreme variations in datasets. They are valid but it makes sense to remove them. You could use standard deviation or domain knowledge to remove them. For example, the sqft per bedroom should have a logical number and if it is out of this number, this should be removed.

As a data scientist when you have a conversation with your business manager (who has expertise in real estate), he will tell you that normally square ft per bedroom is 300 (i.e. 2 bhk apartment is minimum 600 sqft. If you have for example 400 sqft apartment with 2 bhk than that seems suspicious and can be removed as an outlier. We will remove such outliers by keeping our minimum threshold per bhk to be 300 sqft

```
In [40]: df5[df5.total_sqft/df5.bhk<300].head()
```

```
Out[40]:
```

| | location | size | total_sqft | bath | price | bhk | price_per_sqft |
|----|---------------------|-----------|------------|------|-------|-----|----------------|
| 9 | other | 6 Bedroom | 1020.0 | 6.0 | 370.0 | 6 | 36274.509804 |
| 45 | HSR Layout | 8 Bedroom | 600.0 | 9.0 | 200.0 | 8 | 33333.333333 |
| 58 | Murugeshpalya | 6 Bedroom | 1407.0 | 4.0 | 150.0 | 6 | 10660.980810 |
| 68 | Devarachikkanahalli | 8 Bedroom | 1350.0 | 7.0 | 85.0 | 8 | 6296.296296 |
| 70 | other | 3 Bedroom | 500.0 | 3.0 | 100.0 | 3 | 20000.000000 |

Check above data points. We have 6 bhk apartment with 1020 sqft. Another one is 8 bhk and total sqft is 600. These are clear data errors that can be removed safely

```
In [41]: df5.shape
```

```
Out[41]: (13200, 7)
```

```
In [42]: # Remove Outlier for total_sqft
df6 = df5[~(df5.total_sqft/df5.bhk<300)]
df6.shape
```

```
Out[42]: (12456, 7)
```

Outlier Removal Using Standard Deviation and Mean

```
In [43]: df6.price_per_sqft.describe()
```

```
Out[43]: count      12456.000000
mean         6308.502826
std          4168.127339
min           267.829813
25%          4210.526316
50%          5294.117647
75%          6916.666667
max         176470.588235
Name: price_per_sqft, dtype: float64
```

Here we find that min price per sqft is 267 rs/sqft whereas max is 12000000, this shows a wide variation in property prices. We should remove outliers per location using mean and one standard deviation

we are going to write a function that can remove the extreme cases based on std. If our dataset has a normal distribution almost 68% of the dataset is around mean \pm std. So we are going to filter out anything which is beyond std.

```
In [44]: def remove_pps_outliers(df):
          df_out = pd.DataFrame()
          for key, subdf in df.groupby('location'):
              m = np.mean(subdf.price_per_sqft)
              st = np.std(subdf.price_per_sqft)
              reduced_df = subdf[(subdf.price_per_sqft > (m-st)) & (subdf.price_per_sqft < (m+st))]
              df_out = pd.concat([df_out, reduced_df], ignore_index=True)
          return df_out
          df7 = remove_pps_outliers(df6)
          df7.shape
```


```
Out[44]: (10242, 7)
```

With the function we are keeping anything above $m-st$ or below $m+st$ in my reduced dataframe and then we'll keep on appending into my output dataframe.


Let's check if for a given location how does the 2 BHK and 3 BHK property prices look like

```
In [45]: def plot_scatter_chart(df,location):
bhk2 = df[(df.location==location) & (df.bhk==2)]
bhk3 = df[(df.location==location) & (df.bhk==3)]
matplotlib.rcParams['figure.figsize'] = (15,10)
plt.scatter(bhk2.total_sqft,bhk2.price,color='blue',label='2 BHK', s=50)
plt.scatter(bhk3.total_sqft,bhk3.price,marker='+', color='green',label='3 BHK')
plt.xlabel("Total Square Feet Area")
plt.ylabel("Price (Lakh Indian Rupees)")
plt.title(location)
plt.legend()

plot_scatter_chart(df7,"Rajaji Nagar")
```



```
In [46]: plot_scatter_chart(df7, "Hebbal")
```



```
{
  '1' : {
    'mean': 4000,
    'std': 2000,
    'count': 34
  },
  '2' : {
    'mean': 4300,
    'std': 2300,
    'count': 22
  },
}
```

Now we can remove those 2 BHK apartments whose price_per_sqft is less than mean price_per_sqft of 1 BHK apartment


```
In [47]: def remove_bhk_outliers(df):
          exclude_indices = np.array([])
          for location, location_df in df.groupby('location'):
              bhk_stats = {}
              for bhk, bhk_df in location_df.groupby('bhk'):
                  bhk_stats[bhk] = {
                      'mean': np.mean(bhk_df.price_per_sqft),
                      'std': np.std(bhk_df.price_per_sqft),
                      'count': bhk_df.shape[0]
                  }
              for bhk, bhk_df in location_df.groupby('bhk'):
                  stats = bhk_stats.get(bhk-1)
                  if stats and stats['count']>5:
                      exclude_indices = np.append(exclude_indices, bhk_df[bhk_df.price_per_sqft < stats['mean']].index)
          return df.drop(exclude_indices,axis='index')
df8 = remove_bhk_outliers(df7)
# df8 = df7.copy()
df8.shape
```

Out[47]: (7317, 7)


Function Explanation: First I'm doing location groupby going through every location, and for every location dataframe I am creating new dataframe based on bhk. In this new dataframe I am calculating mean, std, and count. Once this for loop is done I am running another for loop trying to exclude those datapoints whose value(price_per_sqft) is less than the mean of the previous bhk.

Plot same scatter chart again to visualize price_per_sqft for 2 BHK and 3 BHK properties

```
In [48]: plot_scatter_chart(df8,"Rajaji Nagar")
```



```
In [49]: plot_scatter_chart(df8, "Hebbal")
```




Cleaning is very great even though there are some data whose not clean completely but that's ok.

Based on above charts we can see that data points highlighted in red below are outliers and they are being removed due to remove_bhk_outliers function

Plot a histogram to see how many apartments I have in per sqft area.


```
In [50]: import matplotlib
matplotlib.rcParams["figure.figsize"] = (20,10)
plt.hist(df8.price_per_sqft,rwidth=0.8)
plt.xlabel("Price Per Square Feet")
plt.ylabel("Count")
```

```
Out[50]: Text(0, 0.5, 'Count')
```



As you see the majority of our data points are between 0 to 10000 price_per_sqft.

Outlier Removal Using Bathrooms Feature


```
In [51]: df8.bath.unique()
```

```
Out[51]: array([ 4.,  3.,  2.,  5.,  8.,  1.,  6.,  7.,  9., 12., 16., 13.])
```

As you see there are some apartment with 12, 13, or 16 bathrooms?!?

```
In [52]: plt.hist(df8.bath,rwidth=0.8)
plt.xlabel("Number of bathrooms")
plt.ylabel("Count")
```

```
Out[52]: Text(0, 0.5, 'Count')
```



```
In [53]: df8[df8.bath>10]
```

```
Out[53]:
```

| | location | size | total_sqft | bath | price | bhk | price_per_sqft |
|-------------|----------------|--------|------------|------|-------|-----|----------------|
| 5277 | Neeladri Nagar | 10 BHK | 4000.0 | 12.0 | 160.0 | 10 | 4000.000000 |
| 8483 | other | 10 BHK | 12000.0 | 12.0 | 525.0 | 10 | 4375.000000 |
| 8572 | other | 16 BHK | 10000.0 | 16.0 | 550.0 | 16 | 5500.000000 |
| 9306 | other | 11 BHK | 6000.0 | 12.0 | 150.0 | 11 | 2500.000000 |
| 9637 | other | 13 BHK | 5425.0 | 13.0 | 275.0 | 13 | 5069.124424 |

Rule of Thump: It is unusual to have 2 more bathrooms than number of bedrooms in a

home

```
In [54]: df8[df8.bath>df8.bhk+2]
```

Out[54]:

| | location | size | total_sqft | bath | price | bhk | price_per_sqft |
|------|---------------|-----------|------------|------|--------|-----|----------------|
| 1626 | Chikkabanavar | 4 Bedroom | 2460.0 | 7.0 | 80.0 | 4 | 3252.032520 |
| 5238 | Nagasandra | 4 Bedroom | 7000.0 | 8.0 | 450.0 | 4 | 6428.571429 |
| 6711 | Thanisandra | 3 BHK | 1806.0 | 6.0 | 116.0 | 3 | 6423.034330 |
| 8408 | other | 6 BHK | 11338.0 | 9.0 | 1000.0 | 6 | 8819.897689 |

Again the business manager has a conversation with you (i.e. a data scientist) that if you have 4 bedroom home and even if you have bathroom in all 4 rooms plus one guest bathroom, you will have total bath = total bed + 1 max. Anything above that is an outlier or a data error and can be removed

```
In [55]: df9 = df8[df8.bath<df8.bhk+2]  
df9.shape
```

Out[55]: (7239, 7)

```
In [56]: df9.head(2)
```

Out[56]:

| | location | size | total_sqft | bath | price | bhk | price_per_sqft |
|---|---------------------|-------|------------|------|-------|-----|----------------|
| 0 | 1st Block Jayanagar | 4 BHK | 2850.0 | 4.0 | 428.0 | 4 | 15017.543860 |
| 1 | 1st Block Jayanagar | 3 BHK | 1630.0 | 3.0 | 194.0 | 3 | 11901.840491 |

Now my dataframe almost clean, so I can start preparing it for machine learning training. For that, I have to drop some unnecessary features. **size** is unnecessary because for size, we already have bhk. **price_per_sqft** can also be dropped because this data is additional and we use it for outlier detection.

```
In [57]: df10 = df9.drop(['size', 'price_per_sqft'],axis='columns')  
df10.head(3)
```

Out[57]:

| | location | total_sqft | bath | price | bhk |
|---|---------------------|------------|------|-------|-----|
| 0 | 1st Block Jayanagar | 2850.0 | 4.0 | 428.0 | 4 |
| 1 | 1st Block Jayanagar | 1630.0 | 3.0 | 194.0 | 3 |
| 2 | 1st Block Jayanagar | 1875.0 | 2.0 | 235.0 | 3 |

Use One Hot Encoding For Location

Machine Learning model cannot interpret text column so we have to convert location into numeric


column. One of the ways to convert a text column into numeric is to use One Hot Encoding.

```
In [58]: dummies = pd.get_dummies(df10.location)
dummies.head(3)
```

Out[58]:

| | 1st Block Jayanagar | 1st Phase JP Nagar | 2nd Phase Judicial Layout | 2nd Stage Nagarbhavi | 5th Block Hbr Layout | 5th Phase JP Nagar | 6th Phase JP Nagar | 7th Phase JP Nagar | 8th Phase JP Nagar | 9th Phase JP Nagar | ... | Vish |
|---|------------------------|-----------------------------|------------------------------------|-------------------------|-------------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|-----|------|
| 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 2 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |

3 rows × 241 columns




```
In [59]: # append dummies into data
# We do not need other data from dummies
df11 = pd.concat([df10,dummies.drop('other',axis='columns')],axis='columns')
df11.head()
```

Out[59]:

| | location | total_sqft | bath | price | bhk | 1st Block Jayanagar | 1st Phase JP Nagar | 2nd Phase Judicial Layout | 2nd Stage Nagarbhavi | 5th Block Hbr Layout | ... | Vij |
|---|------------------------|------------|------|-------|-----|------------------------|-----------------------------|------------------------------------|-------------------------|-------------------------------|-----|-----|
| 0 | 1st Block Jayanagar | 2850.0 | 4.0 | 428.0 | 4 | 1 | 0 | 0 | 0 | 0 | ... | |
| 1 | 1st Block Jayanagar | 1630.0 | 3.0 | 194.0 | 3 | 1 | 0 | 0 | 0 | 0 | ... | |
| 2 | 1st Block Jayanagar | 1875.0 | 2.0 | 235.0 | 3 | 1 | 0 | 0 | 0 | 0 | ... | |
| 3 | 1st Block Jayanagar | 1200.0 | 2.0 | 130.0 | 3 | 1 | 0 | 0 | 0 | 0 | ... | |
| 4 | 1st Block Jayanagar | 1235.0 | 2.0 | 148.0 | 2 | 1 | 0 | 0 | 0 | 0 | ... | |

5 rows × 245 columns



```
In [60]: # Now we can drop location columns because we already covered it with dummies
df12 = df11.drop('location',axis='columns')
df12.head(2)
```

Out[60]:

| | total_sqft | bath | price | bhk | 1st Block Jayanagar | 1st Phase JP Nagar | 2nd Phase Judicial Layout | 2nd Stage Nagarbhavi | 5th Block Hbr Layout | 5th Phase JP Nagar | ... | Vijaya |
|---|------------|------|-------|-----|------------------------|-----------------------------|------------------------------------|-------------------------|-------------------------------|-----------------------------|-----|--------|
| 0 | 2850.0 | 4.0 | 428.0 | 4 | 1 | 0 | 0 | 0 | 0 | 0 | ... | |
| 1 | 1630.0 | 3.0 | 194.0 | 3 | 1 | 0 | 0 | 0 | 0 | 0 | ... | |

2 rows × 244 columns

4. Building a Model

Now our data is clean and is ready for model building. In this part we are going to build a Machine Learning model and then use K fold Cross Validation and gridsearch CV to come up with the best algorithm as well as the best parameters.

```
In [61]: df12.shape
```

Out[61]: (7239, 244)

```
In [62]: X = df12.drop(['price'],axis='columns')
X.head(3)
```

Out[62]:

| | total_sqft | bath | bhk | 1st Block Jayanagar | 1st Phase JP Nagar | 2nd Phase Judicial Layout | 2nd Stage Nagarbhavi | 5th Block Hbr Layout | 5th Phase JP Nagar | 6th Phase JP Nagar | ... | Vijaya |
|---|------------|------|-----|------------------------|-----------------------------|------------------------------------|-------------------------|-------------------------------|-----------------------------|-----------------------------|-----|--------|
| 0 | 2850.0 | 4.0 | 4 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 1 | 1630.0 | 3.0 | 3 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 2 | 1875.0 | 2.0 | 3 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |

3 rows × 243 columns

```
In [63]: X.shape
```

Out[63]: (7239, 243)

```
In [64]: y = df12.price  
y.head(3)
```

```
Out[64]: 0    428.0  
        1    194.0  
        2    235.0  
        Name: price, dtype: float64
```

```
In [65]: len(y)
```

```
Out[65]: 7239
```

train test split

```
In [66]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

```
In [67]: from sklearn.linear_model import LinearRegression  
lr_clf = LinearRegression()  
lr_clf.fit(X_train,y_train)  
lr_clf.score(X_test,y_test)
```

```
Out[67]: 0.8629132245229442
```

Use K Fold cross validation to measure accuracy of LinearRegression model

```
In [68]: from sklearn.model_selection import ShuffleSplit  
from sklearn.model_selection import cross_val_score  
  
cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)  
  
cross_val_score(LinearRegression(), X, y, cv=cv)
```

```
Out[68]: array([0.82702546, 0.86027005, 0.85322178, 0.8436466 , 0.85481502])
```

We can see that in 5 iterations we get a score above 80% all the time. This is pretty good but we want to test few other algorithms for regression to see if we can get even better score. We will use GridSearchCV for this purpose

Find best model using GridSearchCV

```

In [69]: from sklearn.model_selection import GridSearchCV

from sklearn.linear_model import Lasso
from sklearn.tree import DecisionTreeRegressor

def find_best_model_using_gridsearchcv(X,y):
    algos = {
        'linear_regression' : {
            'model': LinearRegression(),
            'params': {
                'normalize': [True, False]
            }
        },
        'lasso': {
            'model': Lasso(),
            'params': {
                'alpha': [1,2],
                'selection': ['random', 'cyclic']
            }
        },
        'decision_tree': {
            'model': DecisionTreeRegressor(),
            'params': {
                'criterion' : ['mse', 'friedman_mse'],
                'splitter': ['best', 'random']
            }
        }
    }
    scores = []
    cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
    for algo_name, config in algos.items():
        gs = GridSearchCV(config['model'], config['params'], cv=cv, return_train_score=False)
        gs.fit(X,y)
        scores.append({
            'model': algo_name,
            'best_score': gs.best_score_,
            'best_params': gs.best_params_
        })

    return pd.DataFrame(scores,columns=['model', 'best_score', 'best_params'])

find_best_model_using_gridsearchcv(X,y)

```

Out[69]:

| | model | best_score | best_params |
|---|-------------------|------------|---|
| 0 | linear_regression | 0.847796 | {'normalize': False} |
| 1 | lasso | 0.726823 | {'alpha': 2, 'selection': 'random'} |
| 2 | decision_tree | 0.717227 | {'criterion': 'friedman_mse', 'splitter': 'best'} |

Based on above results we can say that **LinearRegression** gives the best score. Hence we will use that.

Test the model for few properties

```
In [70]: X.columns
```

```
Out[70]: Index(['total_sqft', 'bath', 'bhk', '1st Block Jayanagar',  
              '1st Phase JP Nagar', '2nd Phase Judicial Layout',  
              '2nd Stage Nagarbhavi', '5th Block Hbr Layout', '5th Phase JP Nagar',  
              '6th Phase JP Nagar',  
              ...  
              'Vijayanagar', 'Vishveshwarya Layout', 'Vishwapriya Layout',  
              'Vittasandra', 'Whitefield', 'Yelachenahalli', 'Yelahanka',  
              'Yelahanka New Town', 'Yelenahalli', 'Yeshwanthpur'],  
              dtype='object', length=243)
```

```
In [71]: np.where(X.columns=='bhk')[0][0]
```

```
Out[71]: 2
```

It returns the location of 'bhk'.

Because Linear Regression is the winner we use this model which we already created `lr_clf` for prediction. Because we have almost 200 locations we have to specify our location for prediction like above example. once you have location index you can set that particular index value to be 1.

```
In [73]: def predict_price(location,sqft,bath,bhk):  
         loc_index = np.where(X.columns==location)[0][0]  
  
         x = np.zeros(len(X.columns))  
         x[0] = sqft  
         x[1] = bath  
         x[2] = bhk  
         if loc_index >= 0:  
             x[loc_index] = 1  
  
         return lr_clf.predict([x])[0]
```

```
In [74]: predict_price('1st Phase JP Nagar',1000, 2, 2)
```

```
Out[74]: 83.86570258312275
```

```
In [75]: predict_price('1st Phase JP Nagar',1000, 3, 3)
```

```
Out[75]: 86.08062284987054
```

```
In [76]: predict_price('Indira Nagar',1000, 2, 2)
```

```
Out[76]: 193.31197733179937
```



```
In [77]: predict_price('Indira Nagar',1000, 3, 3)
```

```
Out[77]: 195.52689759854715
```

Export the tested model to a pickle file

Now you have come this long way with this awesome model and it is ready for production so you need to export it into a pickle file.

```
In [78]: import pickle
with open('bangalore_home_prices_model.pickle','wb') as f:
    pickle.dump(lr_clf,f)
```

The file size of pickle file is 1 KB, this is because this pickle file doesn't have the actual data, it has just model, coefficients.

Export column information into a json file

Other than the model we also need the columns information. For example, in my predict_price function I have X.columns. These columns and their index and their structures are important for making a prediction.

```
In [79]: import json
columns = {
    'data_columns' : [col.lower() for col in X.columns]
}
with open("columns.json","w") as f:
    f.write(json.dumps(columns))
```

| Date | Author |
|------|--------|
|------|--------|

| | |
|------------|---------------------------|
| 11-21-2021 | Ehsan Zia |
|------------|---------------------------|